



cgroups

What is it?

In the Linux kernel, **Control Groups (cgroups)** extend the system by providing additional structures that allow tasks (i.e., processes and threads) to be associated with a specific set of resource management rules (such as CPU limits, memory limits, etc.).

Two principle components:

- A mechanism for **hierarchically grouping** process
- A set of **controllers** (subsystems, kernel components) that manage, control, or monitor processes in cgroups

Hierarchy

- cgroups are organized in a tree-like structure, where each node (cgroup) represents a set of processes.
- Hierarchies allow the creation of parent and child cgroups, where limits can be inherited.

Controllers

Each cgroup can control specific system resources through controllers:

- **CPU (cpu, cpuacct)**: Limit and track CPU usage.
- **Memory (memory)**: Limit and monitor memory usage.
- **Block I/O (blkio)**: Control and monitor disk I/O.
- **Network (net_cls, net_prio)**: Manage network bandwidth and traffic priority.
- Others: **PIDs, devices, freezer**, etc.

A task can be part of **multiple cgroups**, but **only one cgroup per controller** (e.g., CPU, memory, I/O, etc.).

Internals

Tasks and `css_set` (cgroup subsystem state set):

- Each task (process or thread) in the system has a pointer to a `css_set`, which tracks its membership in different cgroups (one per controller).
- `css_set` is a structure that holds pointers to the resource states for each cgroup subsystem (CPU, memory, etc.).

Reference counting:

- The `css_set` is reference-counted. When tasks share the same cgroup membership, they share the same `css_set`.
- Reference counting ensures efficient memory management. When no tasks reference a `css_set`, it can be safely destroyed.

Changing cgroups:

- When a task moves to a new cgroup (e.g., from one memory cgroup to another), the kernel updates the task's `css_set`, creating a new one if needed, and adjusting reference counts.

```
struct task_struct {  
    /* Other fields for process information */  
    struct css_set *cgroups; // Points to the css_set of the task  
};
```

```
struct css_set {  
    struct list_head tasks; // List of tasks in this css_set  
    struct cgroup_subsys_state *subsys[CGROUP_SUBSYS_COUNT]; // Array of subsys states  
    atomic_t refcount; // Reference count for memory management  
};
```

Example

Cgroups: A for CPU, X for Memory, Y for Memory

Task1 and Task2 are both in cgroup A and cgroup X
Move Task1 to Y cgroup

```
shared_css_set.refcount--;
```

```
// new css_set
```

```
struct css_set new_css_set_task1 = {  
    .subsys[CPU_SUBSYS] = cgroup_A_cpu_subsys_state, // Same CPU cgroup (cgroup A)  
    .subsys[MEMORY_SUBSYS] = cgroup_Y_mem_subsys_state, // New memory cgroup (cgroup Y)  
    .refcount = 1, // New reference count for this css_set  
};
```

```
task1.cgroups = &new_css_set_task1;
```

```
struct css_set shared_css_set = {  
    .subsys[CPU_SUBSYS] = cgroup_A_cpu_subsys_state,
```

```
// Points to cgroup A for CPU  
    .subsys[MEMORY_SUBSYS] =  
    cgroup_X_mem_subsys_state,  
    // Points to cgroup X for memory
```

```
    .refcount = 2 // Both Task 1 and Task 2 reference this  
    css_set };
```

Virtual File System Interaction

- **cgroup filesystem (`cgroupfs`):**
 - cgroups expose their control interfaces via a **virtual filesystem** mounted at `/sys/fs/cgroup/`.
 - Each controller (e.g., CPU, memory, blkio) has its own directory or file under `/sys/fs/cgroup/`, allowing direct interaction with the cgroup hierarchy and subsystems.
- **Key directories and files:**
 - **`/sys/fs/cgroup/`:**
The top-level directory for cgroups, containing subdirectories for different controllers (e.g., `cpu`, `memory`, `blkio`).
 - Inside each controller's directory, there are files that represent control parameters (e.g., `cpu.shares`, `memory.limit_in_bytes`).

```
echo 512M > /sys/fs/cgroup/memory/my_cgroup/memory.limit_in_bytes
```


cgroups v1 vs. cgroups v2

cgroups v1:

- Multiple hierarchies, each for one or more controllers.
- Each controller operates independently, which can lead to inconsistencies in resource management.

cgroups v2:

- Unified hierarchy: All controllers must be attached to a single hierarchy.
- Simplified and more consistent resource control across subsystems.
- Provides better isolation, particularly for containerized environments (used by systemd and container runtimes like Docker).

Resources

<https://www.man7.org/>

The Linux Programming Interface - Michael Kerrisk